

Phone Microbiome/Biosensor Project

James Meadow (jfmeadow at gmail dot com)

This analysis document accompanies a manuscript that reports analysis of microbial communities sampled from smartphone touchscreens, as well as the index fingers and thumbs of their owners. This was a project that was conducted as an educational workshop meant to explore innovative ways to monitor health. The manuscript is currently in review. This document was produced with the [knitr package](#) in R, and all source code can be found on GitHub: https://github.com/jfmeadow/Meadow_etal_Phones

Getting data into shape

The first step is to set a random seed (so results are locked in with those reported in the manuscript) and load some necessary R packages and functions.

```
set.seed(42)
options(scipen=7) # curtail scientific notation
library(vegan)
```

```
Loading required package: permute
Loading required package: lattice
This is vegan 2.0-10
```

```
library(labdsv)
```

```
Loading required package: mgcv
Loading required package: nlme
This is mgcv 1.7-28. For overview type 'help("mgcv-package")'.
Loading required package: MASS
```

```
Attaching package: 'labdsv'
```

```
The following object is masked from 'package:stats':
```

```
density
```

```
library(miscTools)
library(xtable)
library(boot)
```

```
Attaching package: 'boot'
```

```
The following object is masked from 'package:lattice':
```

```
melanoma
```

```
# These options just for debugging - knitr is automatically in dir.
# setwd('~/.Dropbox/rwjf/Meadow_et al_Phones/')
# load('phones.RData')

source('functions.R')
```

The OTU table is brought in with a custom function `QiimeIn` that reads a classic OTU table and then splits it into a few useful pieces in a big list. The Bioconductor package `phyloseq` actually has more efficient ways to do this with the `.biom` format, but my modest function will work here - the outcome is practically the same in this case.

```
rw.list <- QiimeIn(file='phones_otu_table.txt')
# removed comment character from first line so R takes it.
rw.map <- read.delim('phones_map.txt', head=TRUE, row.names=1)
rw.big <- rw.list$Table
rw.taxo <- rw.list$Taxa
rm(rw.list)
```

Next, the OTU table needs to be put into shape. R wants to see letters, not numbers, as row names, so a big X is inserted. Remove this. We found a total of 3207836 sequences, consisting of 56 samples and 34400 OTUs, defined at 97% sequence similarity. Then line up with the order of the mapping file.

```
row.names(rw.big) <- gsub('X', '', row.names(rw.big))
rw.big <- rw.big[row.names(rw.map), ]
```

Configure OTU table

Before anything else, we should remove any OTUs from lab controls that showed up in experiment samples. Cell phone surfaces hold really low biomass, so amplification contamination is inevitable. Best to just remove them all. Save the list of contaminant OTUs to shore up the taxonomy table below.

```
cont <- grep('cont', row.names(rw.map))
cont.table <- rw.big[cont, ]
cont.otus <- which(colSums(cont.table) > 0)

rw.table.nocontrol <- rw.big[-cont, -cont.otus]
sort(rowSums(rw.table.nocontrol))
```

```
20.phone 22.phone 31.phone 29.phone 30.phone 35.index 26.thumb 35.phone
      2719      4142      4612       5425       6356       6863       7952      8286
20.thumb 20.index 28.index 19.phone 18.index 18.phone 25.phone 26.phone
      9362     10023     10417     10834     10972     11022     11800     12105
17.phone 34.phone 29.index 30.index 30.thumb 35.thumb 18.thumb 17.thumb
     12348     12730     13474     13898     14052     14241     14286     14833
34.thumb 29.thumb 34.index 28.phone 32.index 22.index 33.thumb 25.thumb
     15015     15392     16022     16946     17051     17250     17290     17543
33.phone 25.index 26.index 33.index 17.index 23.phone 19.index 32.thumb
     18673     18991     19741     19751     20257     20379     20920     20949
32.phone 31.index 23.thumb 22.thumb 19.thumb 23.index 31.thumb 24.phone
```

```

20976    21406    21747    25087    27009    27196    29835    30507
28.thumb 24.thumb 24.index
31828    32916    34054

```

```
rm(rw.big, cont.table, cont)
```

Next:

- Plant sequences are always in 16S datasets, so one way to remove them is to call them by name “*Streptophyta*.” These get removed.
- Do the same for mitochondrial sequences.
- Remove lab contaminants identified above.
- Remove OTUs that are represented by only 1 or 2 sequences - these lend little to community analysis and slow down the whole works.
- The last step is to rarefy all samples to an even sampling depth, in this case 2500 sequences per sample.

```

rw.taxo <- rw.taxo[-cont.otus, ]
streptophyta <- grep('Streptophyta', rw.taxo$taxa.names)
mitochondria <- grep('mitochondria', rw.taxo$taxa.names)
rw.table.tmp <- rw.table.nocontrol[, -c(streptophyta, mitochondria)]
sort(rowSums(rw.table.tmp))

```

```

20.phone 22.phone 31.phone 30.phone 29.phone 35.phone 30.thumb 20.thumb
2660     4005     4471     4731     5306     5633     5993     6268
35.index 30.index 26.thumb 28.index 20.index 18.phone 17.phone 18.index
6572     6705     6719     6758     8873     9711     10212    10669
19.phone 25.phone 26.phone 34.phone 29.index 18.thumb 35.thumb 17.thumb
10825    11591    12067    12651    13196    13832    14059    14765
34.thumb 29.thumb 34.index 28.phone 22.index 32.index 33.thumb 25.thumb
14940    15242    15963    16790    16890    16977    17070    17434
33.phone 26.index 25.index 33.index 17.index 23.phone 19.index 32.thumb
18562    18885    18898    19642    20076    20208    20845    20879
32.phone 31.index 23.thumb 24.phone 28.thumb 22.thumb 19.thumb 23.index
20960    20977    21706    22336    24423    24887    26952    27103
31.thumb 24.thumb 24.index
29752    32669    33865

```

```
rw.table.tmp <- rw.table.tmp[, -c(which(colSums(rw.table.tmp) < 3))]
```

```

rw.25 <- rrarefy(rw.table.tmp, 2500)
rm(streptophyta, mitochondria, rw.table.tmp, cont.otus, rw.table.nocontrol)

```

Since lots of OTUs were removed from the OTU table, we remove them from the taxonomy table - we will want everything lined up downstream.

```

# taxonomy
rw.taxo.25 <- rw.taxo[colnames(rw.25), ]
rm(rw.taxo)

```

So we're left with 127500 sequences in 51 samples and 6667 OTUs.

Configure sample metadata

The mapping file (metadata for each sample) was loaded in during the first step. First, we line up samples with the OTU table row names since it is now in shape. Then there is lots of baggage that comes along with mapping files. Factor variables must be retrained, and then we add three colors that will be used in analysis.

```
# mapping file
map <- rw.map[row.names(rw.25), ] # reorder to match
rm(rw.map) # remove old one

# then reorder a few factors for convenience.
map$individ <- factor(map$individ, levels=c(as.character(levels(map$individ)[1:17])))
map$location <- factor(map$location, levels=c('index', 'thumb', 'phone'))
map$type <- factor(map$type, levels=c('c', 'o', 'p'))
map$dominance <- factor(map$dominance, levels=c('r', 'l'))
map$gender <- factor(map$gender, levels=c('f', 'm'))
map$wash <- factor(map$wash, levels=c('y', 'n'))
map <- map[, c(3, 5, 7, 8, 9)]

# create colors for plotting ease
map$bg <- 'gray30' # phones
map$bg[map$location == 'index'] <- 'cornflowerblue'
map$bg[map$location == 'thumb'] <- 'darkorange'
```

And create a few more variables.

```
map$loc.gen <- paste(map$location, map$gender, sep='.')
map$location2 <- as.character(map$location)
map$location2[map$location2 != 'phone'] <- 'finger'
map$location2 <- factor(map$location2)
map$loc.gen2 <- factor(paste(map$location2, map$gender, sep='.'),
                      levels=c('finger.f', 'phone.f', 'finger.m', 'phone.m'))

index <- which(map$location == 'index')
thumb <- which(map$location == 'thumb')
finger <- which(map$location2 == 'finger')

m <- which(map$gender == 'm')
f <- which(map$gender == 'f')
p <- which(map$location == 'phone')
```

Also, a quick function to print simple L^AT_EX~tables over and over.

```
phtable <- function(tab, capt='') {
  print(xtable(tab, caption=capt), comment=FALSE, timestamp=FALSE)#, type='html')
}
```

Generate taxonomy figure

Taxonomy information, as QIIME gives it, is pretty useless raw. So we have to parse this into a workable data frame, and then use that for figures. First, rename and save on typing! Then the separation between taxonomic levels is used to split strings. A couple more steps and then we have a data frame with 7 taxonomic levels and one last column for total abundance across the rarefied dataset.

```

tt <- rw.taxo.25
tt2 <- as.character(gsub('[:alpha:]{1,1}\\_\\_\\_', '', tt$taxa.names))
tt3 <- strsplit(tt2, split='; ')
ttl <- unlist(lapply(tt3, length))

tt4 <- data.frame(
  kingdom=sapply(tt3, function(x){x[1]}),
  phylum=sapply(tt3, function(x){x[2]}),
  class=sapply(tt3, function(x){x[3]}),
  order=sapply(tt3, function(x){x[4]}),
  family=sapply(tt3, function(x){x[5]}),
  genus=sapply(tt3, function(x){x[6]}),
  species=sapply(tt3, function(x){x[7]}))

tt4$kingdom <- as.character(tt4$kingdom)
tt4$phylum <- as.character(tt4$phylum)
tt4$class <- as.character(tt4$class)
tt4$order <- as.character(tt4$order)
tt4$family <- as.character(tt4$family)
tt4$genus <- as.character(tt4$genus)
tt4$species <- as.character(tt4$species)

for (i in 1:ncol(tt4)){
  tt4[which(is.na(tt4[, i])), i] <- ''
} # warning suppressed

taxo <- tt4
taxo$abundance <- colSums(rw.25)
row.names(taxo) <- rw.taxo.25$qiime.id
rm(tt, tt2, tt3, ttl, tt4)

head(taxo)

```

	kingdom	phylum	class	order
3	Bacteria	Proteobacteria	Gammaproteobacteria	Pseudomonadales
11	Bacteria	Proteobacteria	Gammaproteobacteria	Enterobacteriales
16	Bacteria	Tenericutes	Mollicutes	Acholeplasmatales
30	Bacteria	Planctomycetes	Planctomycetia	Gemmatales
31	Bacteria	Firmicutes	Clostridia	Clostridiales
42	Bacteria	Chloroflexi	C0119	
	family	genus	species	abundance
3	Moraxellaceae	Enhydrobacter	aerosaccus	2
11	Enterobacteriaceae			0
16	Acholeplasmataceae	Acholeplasma		0
30	Isosphaeraceae			3
31	Veillonellaceae	Dialister		226
42				3

Looks good. Then we want to know about the most abundant phyla, to be used for a taxonomy figure. But first, what was the most abundant thing in the dataset?

```
medAbu5 <- sort(apply(rw.25, 2, median)/25, TRUE)[1:5]
medAbu5.m <- sort(apply(rw.25[m, ], 2, median)/25, TRUE)[1:5]
medAbu5.f <- sort(apply(rw.25[f, ], 2, median)/25, TRUE)[1:5]
medAbu5.p <- sort(apply(rw.25[p, ], 2, median)/25, TRUE)[1:5]
medAbu5.fing <- sort(apply(rw.25[finger, ], 2, median)/25, TRUE)[1:5]

phtable(data.frame(taxo[names(medAbu5), -c(1,7,8)], medAbu=medAbu5),
  capt='Most median abundant overall.')
```

	phylum	class	order	family	genus	medAbu
2485	Firmicutes	Bacilli	Gemellales	Gemellaceae	Gemella	2.64
25857	Firmicutes	Bacilli	Lactobacillales	Carnobacteriaceae	Granulicatella	2.04
33170	Actinobacteria	Actinobacteria	Actinomycetales	Micrococcaceae	Rothia	1.80
25456	Actinobacteria	Actinobacteria	Actinomycetales	Actinomycetaceae	Actinomyces	1.68
260	Firmicutes	Bacilli	Lactobacillales	Streptococcaceae	Lactococcus	1.24

Table 1: Most median abundant overall.

```
phtable(data.frame(taxo[names(medAbu5.m), -c(1,7,8)], medAbu=medAbu5.m),
  capt='Most median abundant for men.')
```

	phylum	class	order	family	genus	medAbu
2485	Firmicutes	Bacilli	Gemellales	Gemellaceae	Gemella	1.60
25857	Firmicutes	Bacilli	Lactobacillales	Carnobacteriaceae	Granulicatella	1.28
260	Firmicutes	Bacilli	Lactobacillales	Streptococcaceae	Lactococcus	1.24
25456	Actinobacteria	Actinobacteria	Actinomycetales	Actinomycetaceae	Actinomyces	1.16
33170	Actinobacteria	Actinobacteria	Actinomycetales	Micrococcaceae	Rothia	1.08

Table 2: Most median abundant for men.

```
phtable(data.frame(taxo[names(medAbu5.f), -c(1,7,8)], medAbu=medAbu5.f),
  capt='Most median abundant for women.')
```

	phylum	class	order	family	genus	medAbu
2485	Firmicutes	Bacilli	Gemellales	Gemellaceae	Gemella	3.26
25857	Firmicutes	Bacilli	Lactobacillales	Carnobacteriaceae	Granulicatella	3.08
25456	Actinobacteria	Actinobacteria	Actinomycetales	Actinomycetaceae	Actinomyces	2.90
33170	Actinobacteria	Actinobacteria	Actinomycetales	Micrococcaceae	Rothia	2.46
16496	Bacteroidetes	Bacteroidia	Bacteroidales	Prevotellaceae	Prevotella	1.66

Table 3: Most median abundant for women.

```
phtable(data.frame(taxo[names(medAbu5.p), -c(1,7,8)], medAbu=medAbu5.p),
  capt='Most median abundant on phones.')
```

```
phtable(data.frame(taxo[names(medAbu5.fing), -c(1,7,8)], medAbu=medAbu5.fing),
  capt='Most median abundant on fingers.')
```

	phylum	class	order	family	genus	medAbu
2485	Firmicutes	Bacilli	Gemellales	Gemellaceae	Gemella	3.36
25857	Firmicutes	Bacilli	Lactobacillales	Carnobacteriaceae	Granulicatella	2.32
33170	Actinobacteria	Actinobacteria	Actinomycetales	Micrococcaceae	Rothia	2.20
25456	Actinobacteria	Actinobacteria	Actinomycetales	Actinomycetaceae	Actinomyces	1.80
26366	Bacteroidetes	Bacteroidia	Bacteroidales	[Paraprevotellaceae]	[Prevotella]	1.36

Table 4: Most median abundant on phones.

	phylum	class	order	family	genus	medAbu
2485	Firmicutes	Bacilli	Gemellales	Gemellaceae	Gemella	2.34
25857	Firmicutes	Bacilli	Lactobacillales	Carnobacteriaceae	Granulicatella	1.78
33170	Actinobacteria	Actinobacteria	Actinomycetales	Micrococcaceae	Rothia	1.76
260	Firmicutes	Bacilli	Lactobacillales	Streptococcaceae	Lactococcus	1.66
25456	Actinobacteria	Actinobacteria	Actinomycetales	Actinomycetaceae	Actinomyces	1.58

Table 5: Most median abundant on fingers.

For all groups, it is a *Gemella* OTU. The median abundance was used because OTU abundances can be easily skewed by outlier values (as is also the case in this dataset). So although *Lactobacillus* is the most abundant in the figure produced below, it is not the most consistently abundant when median abundance is considered.

```
ph <- aggregate(taxo$abundance, by=list(taxo$phylum), FUN=sum)
ph[rev(order(ph$x))[1:10], ] # cut off conveniently at unidentified.
```

	Group.1	x
15	Firmicutes	49690
21	Proteobacteria	28145
4	Actinobacteria	26953
6	Bacteroidetes	12869
16	Fusobacteria	4124
1		3942
11	Cyanobacteria	384
25	Tenericutes	306
2	[Thermi]	247
3	Acidobacteria	207

We can use the top 5 and group all others. So a new data frame is created to hold the mean abundances grouped by phylum and by location type (index, thumb, or phone). This code is not pretty but it works.

```
ph.mean <- data.frame(
  Firmicutes = aggregate(rowSums(rw.25[, which(taxo$phylum == 'Firmicutes')]),
    by=list(map$location), FUN=mean),
  Proteobacteria = aggregate(rowSums(rw.25[, which(taxo$phylum == 'Proteobacteria')]),
    by=list(map$location), FUN=mean),
  Actinobacteria = aggregate(rowSums(rw.25[, which(taxo$phylum == 'Actinobacteria')]),
    by=list(map$location), FUN=mean),
  Bacteroidetes = aggregate(rowSums(rw.25[, which(taxo$phylum == 'Bacteroidetes')]),
    by=list(map$location), FUN=mean),
  Fusobacteria = aggregate(rowSums(rw.25[, which(taxo$phylum == 'Fusobacteria')]),
    by=list(map$location), FUN=mean),
  Other = aggregate(rowSums(rw.25[, -c(which(taxo$phylum %in%
    c('Firmicutes', 'Proteobacteria', 'Actinobacteria',
    'Fusobacteria', 'Bacteroidetes')))]),
    by=list(map$location), FUN=mean))

ph.mean <- ph.mean[, c(2, 4, 6, 8, 10, 12)]
row.names(ph.mean) <- c('index', 'thumb', 'phone')
names(ph.mean) <- gsub('.x', '', names(ph.mean))
ph.mean <- ph.mean/2500
```

Then the same thing is done, but to generate standard errors for bar graph error bars. It is the same big ugly code chunk, but FUN=sd is used as the final argument. SE must be calculated by hand in R, so there is one extra step, and then they are reversed so the big bars are on top in descending order.

```
se <- function(x) {sd(x)/sqrt(length(x))}
rw.25.rel <- rw.25/2500

ph.se <- data.frame(
  Firmicutes = aggregate(rowSums(rw.25.rel[, which(taxo$phylum == 'Firmicutes')]),
```



```

by=list(map$location), FUN=se),
Proteobacteria = aggregate(rowSums(rw.25.rel[, which(taxo$phylum == 'Proteobacteria')]),
by=list(map$location), FUN=se),
Actinobacteria = aggregate(rowSums(rw.25.rel[, which(taxo$phylum == 'Actinobacteria')]),
by=list(map$location), FUN=se),
Bacteroidetes = aggregate(rowSums(rw.25.rel[, which(taxo$phylum == 'Bacteroidetes')]),
by=list(map$location), FUN=se),
Fusobacteria = aggregate(rowSums(rw.25.rel[, which(taxo$phylum == 'Fusobacteria')]),
by=list(map$location), FUN=se),
Other = aggregate(rowSums(rw.25.rel[, -c(which(taxo$phylum %in%
c('Firmicutes', 'Proteobacteria', 'Actinobacteria',
'Fusobacteria', 'Bacteroidetes'))])),
by=list(map$location), FUN=se))
ph.se <- ph.se[, c(2, 4, 6, 8, 10, 12)]
row.names(ph.se) <- c('index', 'thumb', 'phone')
names(ph.se) <- gsub('.x', '', names(ph.se))

```

And then turn them upside down for nicer plotting.

```

ph.mean <- ph.mean[, c(6:1)]
ph.se <- ph.se[, c(6:1)]

```

Now we have data in place to make a barplot by hand:

```

# pdf('phylumBarplot.pdf', height=6, width=6, useDingbats=FALSE)
par(mar=c(5,7,2,2), las=1, font.lab=1)
mids <- barplot(as.matrix(ph.mean), beside=TRUE, horiz=TRUE, las=1, xlim=c(0,.6),
border='white', xlab='', axisnames=FALSE,
col=c('cornflowerblue', 'darkorange', 'gray30'), xaxt='n', font.lab=2)
abline(v=c(seq(.1, .6, .1)), col='white', lwd=.5)
arrows(unlist(c(ph.mean-ph.se)), unlist(c(mids)),
unlist(c(ph.mean+ph.se)), unlist(c(mids)),
code=3, angle=90, length=.01)
axis(1, at=c(0,.1,.2,.3,.4,.5,.6), labels=c(0,10,20,30,40,50,60))
legend(.40, 7, legend=c('phone', 'thumb', 'index'), pch=15, pt.cex=2,
col=c('gray30', 'darkorange', 'cornflowerblue'), bty='n', y.intersp=.82)
mtext('Percent of Each Sample', side=1, line=2.4, font=2)
mtext(names(ph.mean), side=2, at=c(mids[2, ]), line=.2, font=3)

```

```

# dev.off()

```

Now try to hone in on the Firmicutes since that is the most prominent difference:

Note - One OTU (Paenibacillus) appears to be an outlier with a wacky distribution, so it is left out. Although set.seed is used at the top, it might show up in the top 10 again if it is all run again.

```

#Firmicutes = aggregate(rowSums(rw.25.rel[, which(taxo$phylum == 'Firmicutes')]),
# by=list(map$location), FUN=se),
fir.table <- rw.25.rel[, which(taxo$phylum == 'Firmicutes')]
fir.taxo <- taxo[taxo$phylum == 'Firmicutes', ]
identical(colnames(fir.table), row.names(fir.taxo))

```

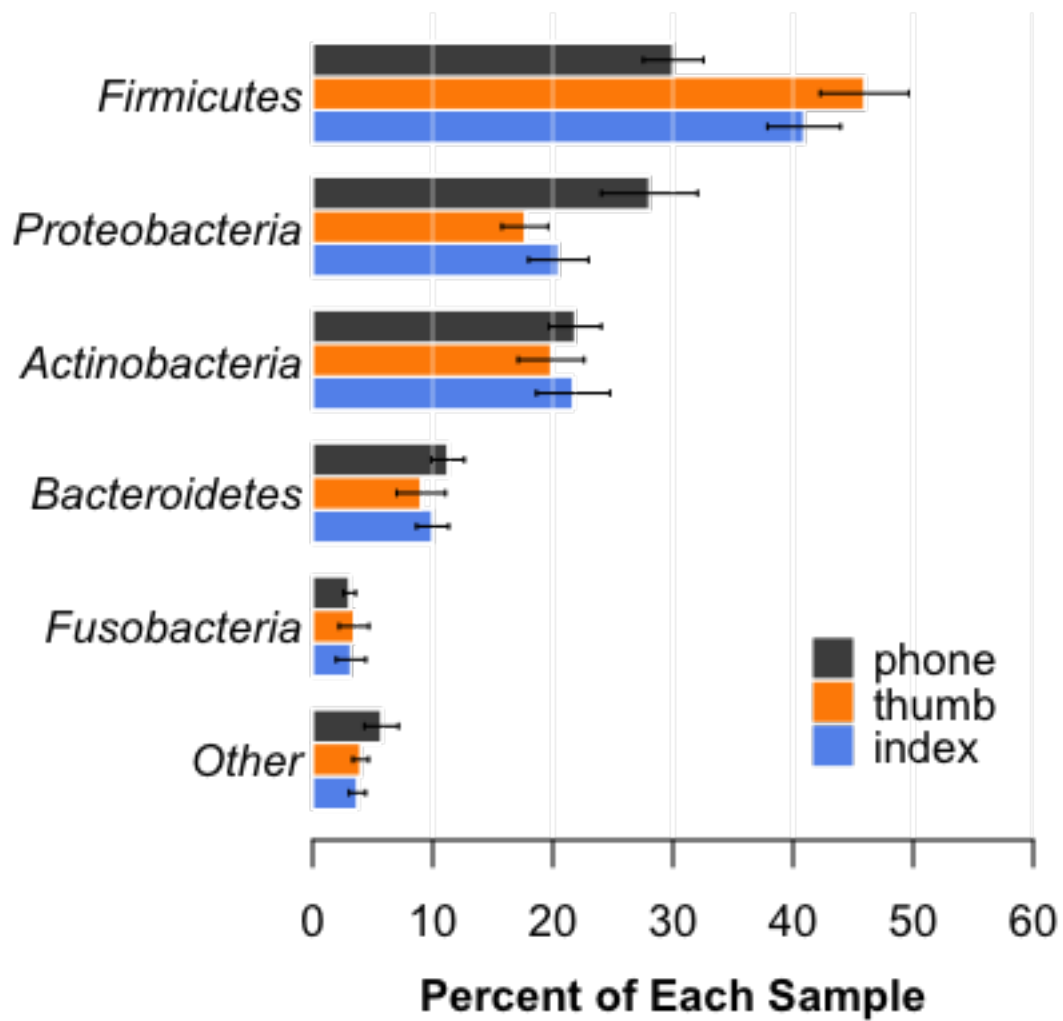


Figure 1: plot of chunk phylumBarplot

```
[1] TRUE
```

```
# leave out outlier
fir.table.10 <- fir.table[, names(rev(sort(colSums(fir.table)))[c(1:3, 5:11)])]
fir.taxo.10 <- fir.taxo[colnames(fir.table.10), ]
dim(fir.table.10)
```

```
[1] 51 10
```

```
sum(colMeans(fir.table.10))
```

```
[1] 0.1557
```

What percentage of observations are represented here? 15.5749%

Look at each of the top 10 to see if there is more to investigate.

Note - This is not evaluated for the knitr document. Only during data exploration.

```
for(i in 1:ncol(fir.table.10)) {
  boxplot(fir.table.10[, i] ~ map$loc.gen2)
  mtext(fir.taxo.10$genus[i])
  readline('press enter')
}
```

Yes. Very much so. There are some really interesting gender differences in the most abundant taxa. This will make a nice addition to the barplot.

Note, again, that one of the original top 10 shows an ugly distribution = all driven by one person with A LOT of Paenibacillus on his fingers. So that one was removed in the code above for clarity - now we're looking at c(1:4, 6:11) to make an even 10. For further investigation, that OTU was ID# 29684. Oddly, it was elevated on both fingers, but didn't show up on his phone. It was nearly absent from all other samples.

Now create data frames to hold plotting info.

```
fir.mean <- data.frame(matrix(NA, 4, 10))
row.names(fir.mean) <- levels(map$loc.gen2)
names(fir.mean) <- colnames(fir.table.10)

fir.se <- data.frame(matrix(NA, 4, 10))
row.names(fir.se) <- levels(map$loc.gen2)
names(fir.se) <- colnames(fir.table.10)

test <- aggregate(fir.table.10[, i], by=list(map$loc.gen2), mean)

if(identical(row.names(fir.mean), as.character(test$Group.1))) {
  for (i in 1:10) {
    fir.mean[, i] <- aggregate(fir.table.10[, i], by=list(map$loc.gen2), mean)$x
    fir.se[, i] <- aggregate(fir.table.10[, i], by=list(map$loc.gen2), se)$x
  }
} else {print("Didn't work - rows weren't lined up!")}

fir.mean <- fir.mean[4:1, 10:1] * 100
fir.se <- fir.se[4:1, 10:1] * 100
```

And do the same for Proteobacteria.

```
pro.table <- rw.25.rel[, which(taxo$phylum == 'Proteobacteria')]
pro.taxo <- taxo[taxo$phylum == 'Proteobacteria', ]
identical(colnames(pro.table), row.names(pro.taxo))
```

```
[1] TRUE
```

```
pro.table.10 <- pro.table[, names(rev(sort(colSums(pro.table)))[c(1:10)])]
pro.taxo.10 <- fir.taxo[colnames(pro.table.10), ]
dim(pro.table.10)
```

```
[1] 51 10
```

```
sum(colMeans(pro.table.10))
```

```
[1] 0.102
```

What percentage of observations are represented here? 10.2016%

```
pro.mean <- data.frame(matrix(NA, 4, 10))
row.names(pro.mean) <- levels(map$loc.gen2)
names(pro.mean) <- colnames(pro.table.10)

pro.se <- data.frame(matrix(NA, 4, 10))
row.names(pro.se) <- levels(map$loc.gen2)
names(pro.se) <- colnames(pro.table.10)

test <- aggregate(pro.table.10[, i], by=list(map$loc.gen2), mean)

if(identical(row.names(pro.mean), as.character(test$Group.1))) {
  for (i in 1:10) {
    pro.mean[, i] <- aggregate(pro.table.10[, i], by=list(map$loc.gen2), mean)$x
    pro.se[, i] <- aggregate(pro.table.10[, i], by=list(map$loc.gen2), se)$x
  }
} else {print("Didn't work - rows weren't lined up!")}

pro.mean <- pro.mean[4:1, 10:1] * 100
pro.se <- pro.se[4:1, 10:1] * 100
```

So now try to combine them. One error bar spills of the right margin, but it is not worth throwing off the whole balance of the figure. So instead I wrote a simple break at the margin to report the real value. *Might move around if re-rarefied without setting seed. So it goes.*

```
# pdf('phylumFermiProteoBarplot.pdf', height=6, width=12, useDingbats=FALSE)
# par(mfrow=c(1,2))
layout(matrix(c(1,2,3), 1,3), widths=c(1.25, 1, 1))
par(mar=c(5,9,2,2), las=1, font.lab=1,
    fg='gray20', col.axis='gray20', col.lab='gray20')
mids <- barplot(as.matrix(ph.mean), beside=TRUE, horiz=TRUE, las=1, xlim=c(0,.6),
```

```

border='white', xlab='', axisnames=FALSE,
col=c('cornflowerblue', 'darkorange', 'gray30'), xaxt='n', font.lab=2)
abline(v=c(seq(.1, .6, .1)), col='white', lwd=.5)
arrows(unlist(c(ph.mean-ph.se)), unlist(c(mids)),
       unlist(c(ph.mean+ph.se)), unlist(c(mids)),
       code=3, angle=90, length=.01)
axis(1, at=c(0,.1,.2,.3,.4,.5,.6), labels=c(0,10,20,30,40,50,60))
legend(.40, 4, legend=c('phone', 'thumb', 'index'), pch=15, pt.cex=3, cex=1.4,
      col=c('gray30', 'darkorange', 'cornflowerblue'), bty='n', y.intersp=.9)
mtext('Percent of Each Sample', side=1, line=2.4, font=2)
mtext('(a) Most abundant phyla', side=3, line=0, font=2, at=0, adj=0)
mtext(names(ph.mean), side=2, at=c(mids[2, ]), line=.2, font=1)
par(xpd=TRUE)
segments(0, c(mids[1, ]-.45), 0, c(mids[3, ]+.45))
par(xpd=FALSE)

# Firmicutes
par(mar=c(5,8,2,2), las=1, font.lab=1, xpd=FALSE,
    fg='gray20', col.axis='gray20', col.lab='gray20')
mids <- barplot(as.matrix(fir.mean), beside=TRUE, horiz=TRUE, las=1, xlim=c(0,8),
  border='white', axisnames=FALSE,
  col=c('gray30', 'cornflowerblue'), font.lab=2)
abline(v=c(seq(1, 8, 1)), col='white', lwd=.5)
arrows(unlist(c(fir.mean-fir.se)), unlist(c(mids)),
       unlist(c(fir.mean+fir.se)), unlist(c(mids)),
       code=3, angle=90, length=.01)
mtext('Percent of Each Sample', side=1, line=2.4, font=2)
mtext('(b) Firmicutes', side=3, line=0, font=2, at=0, adj=0)

par(xpd=TRUE)
segments(0, c(mids[1, ]-.45), 0, c(mids[4, ]+.45))
for (i in 1:10) {
  segments(-.05, mean(mids[, i]), -.5, mean(mids[, i]), col='gray60')
  text(-.25, mean(mids[1:2, i]),
       "\\MA", vfont=c("sans serif symbol", "plain"), font=2, col='gray20', cex=1.5)
  text(-.25, mean(mids[3:4, i]),
       "\\VE", vfont=c("sans serif symbol", "plain"), font=2, col='gray20', cex=1.5)
  barname <- taxo[names(fir.mean)[i], 'genus']
  font <- 3
  if (barname == '') {
    barname <- taxo[names(fir.mean)[i], 'family']
    font <- 1}
  mtext(barname, side=2, at=mean(mids[, i]), line=1.5, font=font)
}

# Proteobacteria
par(mar=c(5,8,2,2.5), las=1, font.lab=1, xpd=FALSE,
    fg='gray20', col.axis='gray20', col.lab='gray20')
mids <- barplot(as.matrix(pro.mean), beside=TRUE, horiz=TRUE, las=1, xlim=c(0,8),
  border='white', axisnames=FALSE,
  col=c('gray30', 'cornflowerblue'), font.lab=2)
abline(v=c(seq(1, 8, 1)), col='white', lwd=.5)
arrows(unlist(c(pro.mean-pro.se)), unlist(c(mids)),

```

```

    unlist(c(pro.mean+pro.se)), unlist(c(mids)),
    code=3, angle=90, length=.01)
mtext('Percent of Each Sample', side=1, line=2.4, font=2)
mtext('(c) Proteobacteria', side=3, line=0, font=2, at=0, adj=0)

par(xpd=TRUE)
segments(0, c(mids[1, ]-.45), 0, c(mids[4, ]+.45))
for (i in 1:10) {
  segments(-.05, mean(mids[, i]), -.5, mean(mids[, i]), col='gray60')
  text(-.25, mean(mids[1:2, i]),
        "\\MA", vfont=c("sans serif symbol", "plain"), font=2, col='gray20', cex=1.5)
  text(-.25, mean(mids[3:4, i]),
        "\\VE", vfont=c("sans serif symbol", "plain"), font=2, col='gray20', cex=1.5)
  barname <- taxo[names(pro.mean)[i], 'genus']
  font <- 3
  if (barname == '') {
    barname <- taxo[names(pro.mean)[i], 'family']
    font <- 1 }
  mtext(barname, side=2, at=mean(mids[, i]), line=1.5, font=font)
  if (barname == 'Photobacterium') {fixit <- i}
}

# Fix Photobacterium error bar break
fixx <- 8
fixy <- mids[1, fixit]
realSE <- round(pro.mean[1, fixit] + pro.se[1, fixit], 1)
par(xpd=TRUE)
segments(fixx-.1, fixy-.5, fixx+.1, fixy+.5)
segments(fixx+.1, fixy-.5, fixx+.3, fixy+.5)
text(fixx+.1, fixy, realSE, pos=4, cex=.7)

```

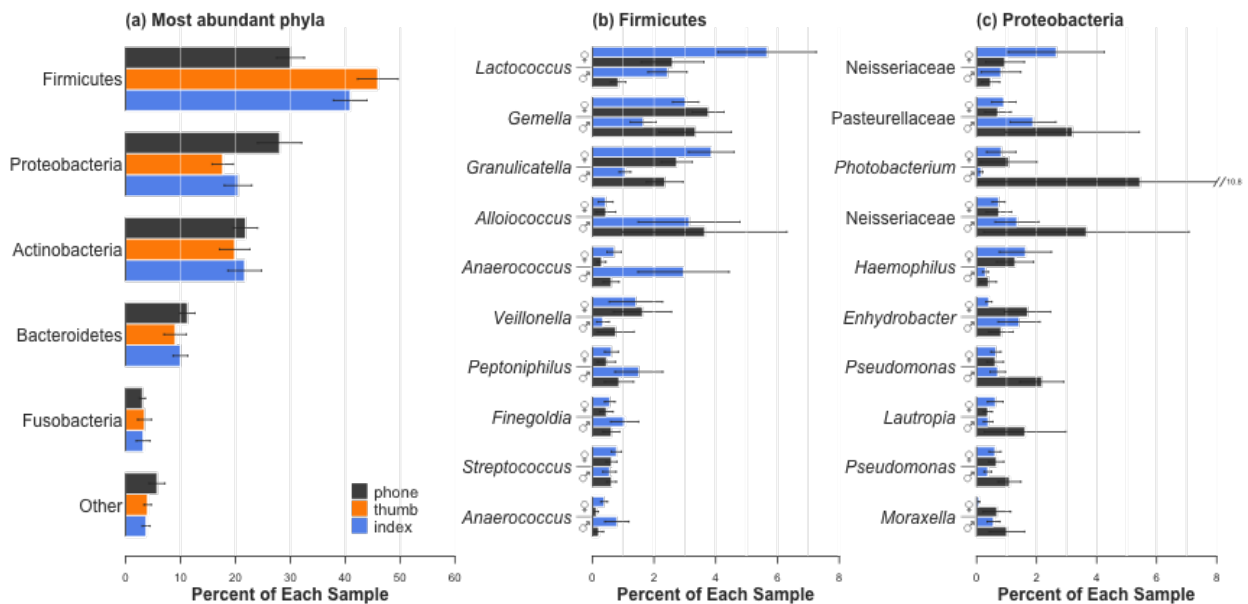


Figure 2: plot of chunk longPhylumBarplots

```
# dev.off()
```

Canberra Distance Barplots - how are phones related to people?

Next, we want to know how communities break out between people and their phones. To do this, we make a distance matrix. In our case, we want to be able to easily explain so we use Jaccard similarity

$$S_{jaccard} = \frac{\text{shared richness}}{\text{combined richness}}$$

so that we can interpret in easy language. Later, we'll also want a *similarity* rather than a *distance*, so we'll invert the distance R gives by default

$$S_{jaccard} = 1 - D_{jaccard}$$

. This way things with more in common have higher values, and that is easier to visualize.

Note that this was tried also with the same Canberra distance that will be used later for ordinations. Results were almost identical, but Jaccard is much easier to interpret for this sort of graph, so we use Jaccard. Also, the Jaccard calculation in the **vegan** package is not the one we want, so I've written one that calculates the original Jaccard similarity. This takes a bit of time since it is not compiled. Thus, this is chached so that it does not have to run every time. If rerunning with slightly different data, uncache this code chunk. So it goes.

```
JacJFM <- function(x=rw.25) {
  mat <- matrix(0, nrow(x), nrow(x))
  mat.names <- dimnames(rw.25)[[1]]
  dimnames(mat) <- list(mat.names, mat.names)

  for(i in 1:nrow(mat)) {
    for(j in 1:ncol(mat)) {
      iname <- row.names(x)[i]
      jname <- row.names(x)[j]
      if(iname != jname) {
        ijcount <- apply(x[c(iname, jname), ], 2, sum)
        ijintersect <- length(which(ijcount == 2))
        iunion <- length(which(ijcount > 0))
        ijvalue <- ijintersect/iunion
        mat[iname, jname] <- ijvalue
      }
    }
  }
  invisible(as.dist(mat))
}

dis <- 1-JacJFM()
```

Since we want to do this several times, I'll package a few tedious routines into functions to cut down on repetitive coding. First set up a data frame for the whole dataset.

```
bar.df <- data.frame(matrix(0,17,3))
names(bar.df) <- c('in.th', 'in.ph', 'th.ph')
row.names(bar.df) <- unique(map$individ)
```

Then create the functions that will be used a few times.

```
makeBarDF <- function() {
  for(i in 1:nrow(bar.df)) {
    bar.df[i, 1] <- as.matrix(dis)[which(map$indiv == row.names(bar.df)[i] &
                                         map$location == 'index'),
                                   which(map$indiv == row.names(bar.df)[i] &
                                         map$location == 'thumb')]]
    bar.df[i, 2] <- as.matrix(dis)[which(map$indiv == row.names(bar.df)[i] &
                                         map$location == 'index'),
                                   which(map$indiv == row.names(bar.df)[i] &
                                         map$location == 'phone')]]
    bar.df[i, 3] <- as.matrix(dis)[which(map$indiv == row.names(bar.df)[i] &
                                         map$location == 'thumb'),
                                   which(map$indiv == row.names(bar.df)[i] &
                                         map$location == 'phone')]]
  }
  invisible(bar.df)
}

makeBarSummary <- function(bar.df=bar.df) {
  bar.summary <- data.frame(cbind(apply(bar.df, 2, mean),
                                   (apply(bar.df, 2, sd)/sqrt(nrow(bar.df))))))
  bar.summary[, 3] <- bar.summary[, 1]-bar.summary[, 2]
  bar.summary[, 4] <- bar.summary[, 1]+bar.summary[, 2]
  names(bar.summary) <- c('mean', 'se', 'se.lo', 'se.hi')
  invisible(bar.summary)
}
```

And then run the data set through the functions.

```
bar.jac.df <- makeBarDF()
bar.jac <- makeBarSummary(bar.df=bar.jac.df)

bar.jac.df
```

	in.th	in.ph	th.ph
17	0.6273	0.7686	0.7886
18	0.6580	0.8397	0.8427
19	0.6426	0.7560	0.7911
20	0.6873	0.8364	0.8830
22	0.6281	0.8246	0.8279
23	0.6564	0.6842	0.7484
24	0.6350	0.8333	0.8149
25	0.6558	0.7968	0.7837
26	0.7771	0.8100	0.8278
28	0.7844	0.7989	0.8305
29	0.7310	0.7896	0.7959
30	0.6292	0.8487	0.8531
31	0.7232	0.8745	0.8597
32	0.7627	0.6952	0.6737
33	0.6221	0.6771	0.7073
34	0.6540	0.7960	0.8017
35	0.8164	0.8543	0.8525


```
bar.jac
```

```
      mean      se se.lo se.hi
in.th 0.6877 0.01563 0.6721 0.7033
in.ph 0.7932 0.01452 0.7787 0.8077
th.ph 0.8049 0.01326 0.7916 0.8181
```

Each additional time, we're only interested in a few samples at a time, so run subsets through the same functions. Each starts out being named generically `bar.df`, but then each object gets put into a uniquely named data frame.

First we need to know how many are in each group.

```
table(map$gender)/3
```

```
  f  m
10  7
```

```
table(map$wash)/3
```

```
 y  n
 9  8
```

```
# females=10
bar.df <- data.frame(matrix(0,10,3))
names(bar.df) <- c('in.th', 'in.ph', 'th.ph')
row.names(bar.df) <- unique(map$indiv[which(map$gender == 'f')])
bar.df.female.j <- makeBarDF()
bar.female.j <- makeBarSummary(bar.df=bar.df.female.j)

# males=7
bar.df <- data.frame(matrix(0,7,3))
names(bar.df) <- c('in.th', 'in.ph', 'th.ph')
row.names(bar.df) <- unique(map$indiv[which(map$gender == 'm')])
bar.df.male.j <- makeBarDF()
bar.male.j <- makeBarSummary(bar.df=bar.df.male.j)

# yes wash=9
bar.df <- data.frame(matrix(0,9,3))
names(bar.df) <- c('in.th', 'in.ph', 'th.ph')
row.names(bar.df) <- unique(map$indiv[which(map$wash == 'y')])
bar.df.wash.j <- makeBarDF()
bar.wash.j <- makeBarSummary(bar.df=bar.df.wash.j)

# no wash=8
bar.df <- data.frame(matrix(0,8,3))
names(bar.df) <- c('in.th', 'in.ph', 'th.ph')
row.names(bar.df) <- unique(map$indiv[which(map$wash == 'n')])
bar.df.nowash.j <- makeBarDF()
bar.nowash.j <- makeBarSummary(bar.df=bar.df.nowash.j)
```

To make it easier to combine side by side barplots, we'll combine them into joined data frames based on their variables (males and females, and wash and no-wash). Then one last step to invert the numbers from a distance to a similarity.

```
bar.mf <- data.frame(rbind(bar.male.j[1, ], bar.female.j[1, ],
                           bar.male.j[2, ], bar.female.j[2, ],
                           bar.male.j[3, ], bar.female.j[3, ]))
bar.wnw <- data.frame(rbind(bar.wash.j[1, ], bar.nowash.j[1, ],
                           bar.wash.j[2, ], bar.nowash.j[2, ],
                           bar.wash.j[3, ], bar.nowash.j[3, ]))

bar.jac <- 1-bar.jac
bar.mf <- 1-bar.mf
bar.wnw <- 1-bar.wnw
```

All data are in place, so there is lots of futzy code to get barplots to look nice. These were modeled after Edward Tufte's *The Visual Display of Quantitative Information*.

```
# pdf('longBarplotFigure.pdf', width=8, height=4)

ylim <- c(0,.45)
layout(matrix(c(1,2,3), 1,3), widths=c(1, 1.6, 1.6))
par(mar=c(4, 4.5, 2, 1), las=1, fg='gray20', lheight=1, col.axis='gray20', col.lab='gray20')
mids <- barplot(bar.jac$mean, las=1,
  border='transparent', axes=FALSE, ylim=ylim, yaxs='i',
  # ylab='Percent of species in common')
  ylab='')
mtext('Jaccard Similarity\n(as % of shared OTUs)', side=2, line=2, las=0, cex=.8)
abline(h=c(.1,.2,.3,.4), col='white', lwd=1)
mtext(c('index\n&\nthumb', 'index\n&\nphone', 'thumb\n&\nphone'),
  side=1, line=2.1, at=c(mids), cex=.7, col='gray20')
axis(2, col='gray20', col.ticks='gray20',
  at=c(0,.1,.2,.3,.4), labels=c(0,10,20,30,'40%'))
arrows(mids, bar.jac$se.lo, mids, bar.jac$se.hi, code=3,
  angle=90, length=.05, col='gray40')
mtext('(a) All samples', font=2, col='gray20', line=.5)
#
par(mar=c(4, 2, 2, 1))
mids <- barplot(bar.wnw$mean, las=1, col=c('gray80', 'gray50'),
  border='transparent', axes=FALSE, ylim=ylim, yaxs='i')
abline(h=c(.1,.2,.3,.4), col='white', lwd=1)
mtext(c('index\n&\nthumb', 'index\n&\nphone', 'thumb\n&\nphone'),
  side=1, line=2.1, cex=.7, col='gray20',
  at=c(mean(mids[1:2,1]), mean(mids[3:4,1]), mean(mids[5:6,1])))
arrows(mids, bar.wnw$se.lo, mids, bar.wnw$se.hi, code=3,
  angle=90, length=.05, col='gray20')
mtext('(b) Did you wash your hands?', font=2, col='gray20', line=.5)
legend(5, .4, legend=c('yes', 'no'), pch=15, col=c('gray80', 'gray50'),
  pt.cex=2, bty='n')
par(xpd=TRUE)
segments(c(mids[c(1,3,5), 1])-.48, rep(-.00, 3),
  c(mids[c(2,4,6), 1])+.48, rep(-.00, 3),
```

```

col='gray30')
par(xpd=FALSE)
#
par(mar=c(4, 2, 2, 1))
mids <- barplot(bar.mf$mean, las=1, col=c('gray80', 'gray50'),
  border='transparent', axes=FALSE, ylim=ylim, yaxs='i')
abline(h=c(.1,.2,.3,.4), col='white', lwd=1)
mtext(c('index\n&\nthumb', 'index\n&\nphone', 'thumb\n&\nphone'),
  side=1, line=2.1, cex=.7, col='gray20',
  at=c(mean(mids[1:2,1]), mean(mids[3:4,1]), mean(mids[5:6,1])))
arrows(mids, bar.mf$se.lo, mids, bar.mf$se.hi, code=3,
  angle=90, length=.05, col='gray20')
mtext('(c) Male or Female?', font=2, col='gray20', line=.5)
legend(5, .4, legend=c('m', 'f'), pch=15, col=c('gray80', 'gray50'),
  pt.cex=2, bty='n')
par(xpd=TRUE)
segments(c(mids[c(1,3,5), 1])-.48, rep(-.00, 3),
  c(mids[c(2,4,6), 1])+.48, rep(-.00, 3),
  col='gray30')

```

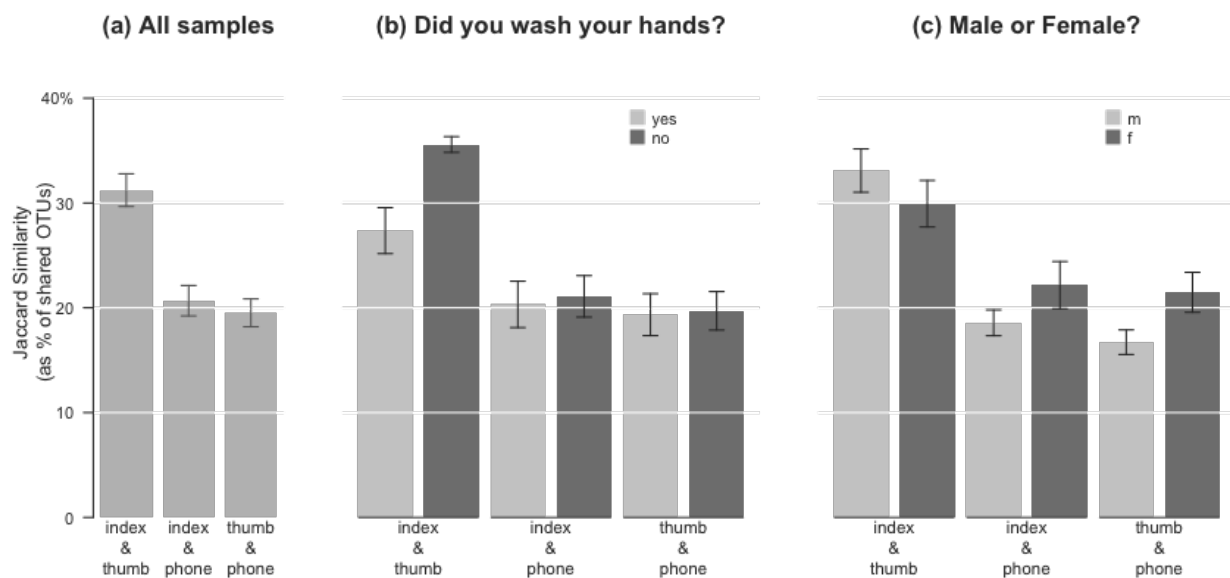


Figure 3: plot of chunk longBarplotFigure

```

par(xpd=FALSE)
#
# dev.off()

```

Isn't that nice? So this figure tells us a couple of really interesting things about the world that we didn't know before! For instance, about 35% of the bacterial taxa we find on our own index finger are also found on the opposing thumb. And (even though fewer) about 20% are also found on our phones! That general pattern is repeated regardless of whether we are looking at men or women, but interestingly, women seem to have more taxa in common with their phones. And your two fingers have more in common if you did not wash your hands.

We can use simple paired t-tests to check some of the patterns in the plots. For instance: Is the first difference significant (are fingers closer to one another than either finger compared to phones)? First, we again need to invert distances to similarities for easier interpretation (already did this in summary tables, but now for raw distance data frames).

```
bar.jac.df <- 1-bar.jac.df
bar.df.male.j <- 1-bar.df.male.j
bar.df.female.j <- 1-bar.df.female.j
bar.df.wash.j <- 1-bar.df.wash.j
bar.df.nowash.j <- 1-bar.df.nowash.j
```

```
t.test(bar.jac.df$in.th, bar.jac.df$in.ph, paired=TRUE)
```

Paired t-test

```
data: bar.jac.df$in.th and bar.jac.df$in.ph
t = 5.435, df = 16, p-value = 0.00005502
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 0.06434 0.14663
sample estimates:
mean of the differences
      0.1055
```

```
t.test(bar.jac.df$in.th, bar.jac.df$th.ph, paired=TRUE)
```

Paired t-test

```
data: bar.jac.df$in.th and bar.jac.df$th.ph
t = 6.106, df = 16, p-value = 0.00001518
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 0.07648 0.15785
sample estimates:
mean of the differences
      0.1172
```

Yes. Very much so. It looks like our fingers have about 35% of their taxa in common, while fingers and phones only share about 20% of taxa.

And how about males and females differentially related to their phones?

```
t.test(bar.df.male.j$in.ph, bar.df.female.j$in.ph)
```

Welch Two Sample t-test

```
data: bar.df.male.j$in.ph and bar.df.female.j$in.ph
t = -1.4, df = 13.4, p-value = 0.1841
alternative hypothesis: true difference in means is not equal to 0
```

```
95 percent confidence interval:
-0.09125  0.01934
sample estimates:
mean of x mean of y
  0.1857   0.2216
```

```
t.test(bar.df.male.j$th.ph, bar.df.female.j$th.ph)
```

Welch Two Sample t-test

```
data: bar.df.male.j$th.ph and bar.df.female.j$th.ph
t = -2.126, df = 14.07, p-value = 0.05172
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
-0.0955972  0.0004063
sample estimates:
mean of x mean of y
  0.1671   0.2147
```

Not so much. And are our fingers more similar if we don't wash our hands?

```
t.test(bar.df.wash.j$in.th, bar.df.nowash.j$in.th)
```

Welch Two Sample t-test

```
data: bar.df.wash.j$in.th and bar.df.nowash.j$in.th
t = -3.545, df = 9.838, p-value = 0.005446
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
-0.13398 -0.03042
sample estimates:
mean of x mean of y
  0.2736   0.3558
```

```
t.test(bar.df.wash.j$in.ph, bar.df.nowash.j$in.ph)
```

Welch Two Sample t-test

```
data: bar.df.wash.j$in.ph and bar.df.nowash.j$in.ph
t = -0.2582, df = 14.97, p-value = 0.7998
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
-0.07089  0.05558
sample estimates:
mean of x mean of y
  0.2032   0.2109
```

The first (two fingers) is suggestive, but not totally convincing - there might be a difference of about 10% of OTUs. The second (relationship to phones) is not different at all.

One last barplot to show whether or not our phones are more indicative of our own microbiome. The workflow will be pretty much the same, but we are picking out:

- the similarities of each person's index finger compared to their own phone, and
- each person's finger compared to the average distance to everyone else's phone

```
dismat <- as.matrix(dis)
dismat <- 1-dismat

bar.df2 <- data.frame(matrix(0,17,2))
names(bar.df2) <- c('same.in.ph', 'others.in.ph')
row.names(bar.df2) <- unique(map$individ)

for(i in 1:nrow(bar.df2)) {
  bar.df2[i, 1] <- dismat[which(map$individ == row.names(bar.df2)[i] &
                               map$location == 'index'),
                        which(map$individ == row.names(bar.df2)[i] &
                               map$location == 'phone')]]
  bar.df2[i, 2] <- mean(dismat[which(map$individ == row.names(bar.df2)[i] &
                                     map$location == 'index'),
                           which(map$individ != row.names(bar.df2)[i] &
                                  map$location == 'phone')]])
}

bar.summary2 <- data.frame(
  cbind(apply(bar.df2, 2, mean),
        (apply(bar.df2, 2, sd)/sqrt(nrow(bar.df2)))))

bar.summary2[, 3] <- bar.summary2[, 1]-bar.summary2[, 2]
bar.summary2[, 4] <- bar.summary2[, 1]+bar.summary2[, 2]
names(bar.summary2) <- c('mean', 'se', 'se.lo', 'se.hi')
bar.others <- bar.summary2
bar.others.df <- bar.df2
```

And make the plot.

```
#pdf('resemblePhone.pdf', width=2.5, height=3)
par(mar=c(3, 4.5, 2, 1), las=1, col.axis='gray20', col.lab='gray20', fg='gray20')
mids <- barplot(bar.others$mean, las=1,
  border='transparent', axes=FALSE, ylim=c(0,.3), yaxs='i',
  # ylab='Percent of species in common')
  ylab='')
mtext('Jaccard Similarity\n(as % of shared OTUs)', side=2, line=2.2, las=0)
abline(h=seq(0, .3, .05), col='white', lwd=1)
mtext(c('index &\nown\nphone', 'index &\nother\nphones'),
  side=1, line=1.6, at=c(mids), cex=.8, col='gray20')
axis(2, col='gray20', col.ticks='gray20',
  at=c(0,.1,.2,.3), labels=c(0,10,20,'30%'))
arrows(mids, bar.others$se.lo, mids, bar.others$se.hi, code=3,
  angle=90, length=.05, col='gray40')
mtext('Does your phone\nresemble you?', font=2, col='gray20')
```

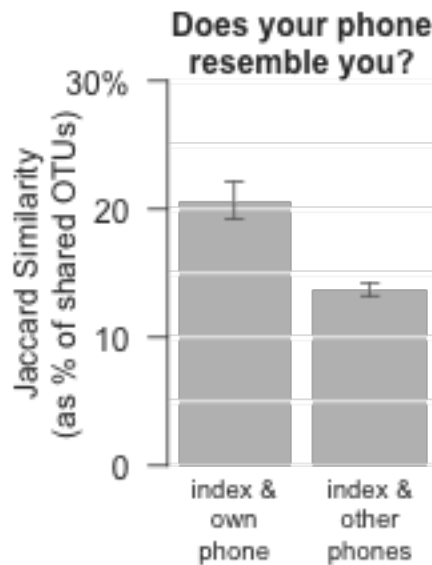


Figure 4: plot of chunk makePersonToPhoneBarplot

```
#dev.off()
```

And is that significant?

```
apply(bar.others.df, 2, mean)
```

```
same.in.ph others.in.ph
0.2068      0.1373
```

```
t.test(bar.others.df$same.in.ph, bar.others.df$others.in.ph, paired=TRUE)
```

Paired t-test

```
data: bar.others.df$same.in.ph and bar.others.df$others.in.ph
t = 4.944, df = 16, p-value = 0.0001466
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 0.03969 0.09929
sample estimates:
mean of the differences
      0.06949
```

So yes, your phone might be able to identify you. Or in other words, we see some evidence that the microbial assemblages on our phones are perhaps extensions of our own, and that they are to some degree personalized to us!

But this distance might just be driven by women, so quick check to see if this result varies by gender. It is lower power, so can't really be compared directly to the test above, but gives us an idea.

```

bar.others.f.df <- bar.others.df[which(map[index, 'gender'] == 'f'), ]

bar.summary.f <- data.frame(
  cbind(apply(bar.others.f.df, 2, mean),
    (apply(bar.others.f.df, 2, sd)/sqrt(nrow(bar.others.f.df))))))

bar.summary.f[, 3] <- bar.summary.f[, 1]-bar.summary.f[, 2]
bar.summary.f[, 4] <- bar.summary.f[, 1]+bar.summary.f[, 2]
names(bar.summary.f) <- c('mean', 'se', 'se.lo', 'se.hi')

bar.others.m.df <- bar.others.df[which(map[index, 'gender'] == 'm'), ]

bar.summary.m <- data.frame(
  cbind(apply(bar.others.m.df, 2, mean),
    (apply(bar.others.m.df, 2, sd)/sqrt(nrow(bar.others.m.df))))))

bar.summary.m[, 3] <- bar.summary.m[, 1]-bar.summary.m[, 2]
bar.summary.m[, 4] <- bar.summary.m[, 1]+bar.summary.m[, 2]
names(bar.summary.m) <- c('mean', 'se', 'se.lo', 'se.hi')

```

And make the plots.

```

# pdf('resemblePhone.pdf', width=2.5, height=3)
par(mar=c(3, 4.5, 2, 1), las=1, col.axis='gray20', col.lab='gray20', fg='gray20')
mids <- barplot(bar.summary.f$mean, las=1,
  border='transparent', axes=FALSE, ylim=c(0,.3), yaxs='i',
  ylab='')
mtext('Jaccard Similarity\n(as % of shared OTUs)', side=2, line=2.2, las=0)
abline(h=seq(0, .3, .05), col='white', lwd=1)
mtext(c('index &\nown\nphone', 'index &\nother\nphones'),
  side=1, line=1.6, at=c(mids), cex=.8, col='gray20')
axis(2, col='gray20', col.ticks='gray20',
  at=c(0,.1,.2,.3), labels=c(0,10,20,'30%'))
arrows(mids, bar.summary.f$se.lo, mids, bar.summary.f$se.hi, code=3,
  angle=90, length=.05, col='gray40')
mtext("Do women resemble\ntheir phones?", font=2, col='gray20')

```

```

par(mar=c(3, 4.5, 2, 1), las=1, col.axis='gray20', col.lab='gray20', fg='gray20')
mids <- barplot(bar.summary.m$mean, las=1,
  border='transparent', axes=FALSE, ylim=c(0,.3), yaxs='i',
  ylab='')
mtext('Jaccard Similarity\n(as % of shared OTUs)', side=2, line=2.2, las=0)
abline(h=seq(0, .3, .05), col='white', lwd=1)
mtext(c('index &\nown\nphone', 'index &\nother\nphones'),
  side=1, line=1.6, at=c(mids), cex=.8, col='gray20')
axis(2, col='gray20', col.ticks='gray20',
  at=c(0,.1,.2,.3), labels=c(0,10,20,'30%'))
arrows(mids, bar.summary.m$se.lo, mids, bar.summary.m$se.hi, code=3,
  angle=90, length=.05, col='gray40')
mtext("Do men resemble\ntheir phones?", font=2, col='gray20')

```

And is that significant for both groups?

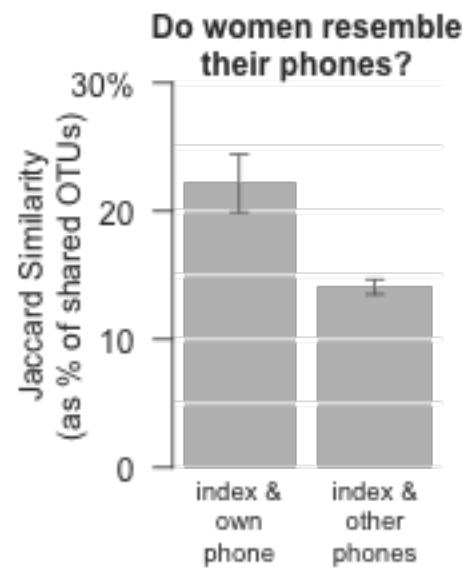


Figure 5: plot of chunk makeGenderToPhoneBarplot

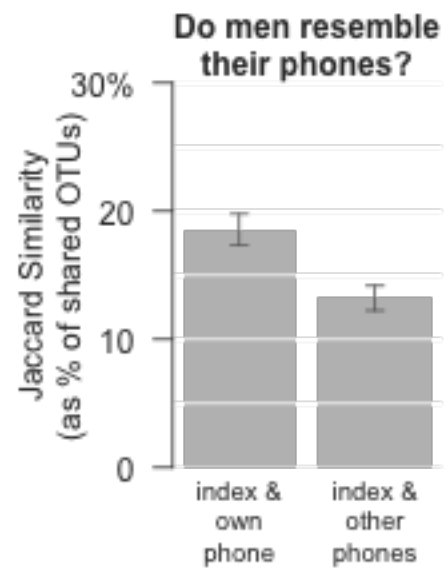


Figure 6: plot of chunk makeGenderToPhoneBarplot

```
diff.f <- round(diff(c(apply(bar.others.f.df, 2, mean))) * (-100), digits=2)
t.test(bar.others.f.df$same.in.ph, bar.others.f.df$others.in.ph, paired=TRUE)
```

Paired t-test

```
data: bar.others.f.df$same.in.ph and bar.others.f.df$others.in.ph
t = 3.691, df = 9, p-value = 0.004993
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 0.03123 0.13017
sample estimates:
mean of the differences
          0.0807
```

```
diff.m <- round(diff(c(apply(bar.others.m.df, 2, mean))) * (-100), digits=2)
t.test(bar.others.m.df$same.in.ph, bar.others.m.df$others.in.ph, paired=TRUE)
```

Paired t-test

```
data: bar.others.m.df$same.in.ph and bar.others.m.df$others.in.ph
t = 3.96, df = 6, p-value = 0.007452
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 0.02044 0.08653
sample estimates:
mean of the differences
          0.05348
```

Yes - both sexes share, on average, more with their own phones than with everyone else's. Men share about 5.35% percent more taxa with their own than with other phones, while women share 8.07% more than they share with other people's phones.

Community differences - Ordination and Discriminant Analysis

Do the fingers of men and women harbor different types of bacteria? Previous research says yes. In the current study, some people washed hands and some didn't. So we should find out if we have a balanced study (i.e., relatively even numbers in all four categories?). The answer is, of course, somewhat funny, though not *significantly* funny.

```
gender.wash <- table(map$gender, map$wash)
cst <- chisq.test(gender.wash)
cst
```

Pearson's Chi-squared test with Yates' continuity correction

```
data: gender.wash
X-squared = 0.8503, df = 1, p-value = 0.3565
```

```
cst$observed
```

```
      y  n  
f 18 12  
m  9 12
```

```
cst$expected
```

```
      y      n  
f 15.88 14.118  
m 11.12  9.882
```

```
rm(cst, gender.wash)
```

Anyway, the sample is reasonably well balanced. PERMANOVA tests like `adonis` are not necessarily robust to big imbalances, but probably not a problem for us.

We'll use the Canberra distance, since we expect most of the abundant taxa to overlap - we are interested in differences among the relatively rare OTUs. Compared to the easily interpretable Jaccard index used above, these sophisticated dissimilarities tend to be more mathematically satisfying for ordination. If we look at an NMDS of all samples, it looks like there is a separation between fingers and phones, but not between finger types. Not too surprising given what we saw in the bar plots above. Notice that the ordination gets put into a generic object `n` to save on typing but also to make it easier to switch distance matrices or `nmDS` objects later.

```
rw.25.can <- vegdist(rw.25, 'canberra')
```

```
rw.25.nmDS.can <- bestnmDS(rw.25.can, k=2)
```

If we emphasize gender, it seems that men and women fall in different parts of the plot, and their communities, when considering all samples, are very significantly different.

```
rw.25.nmDS.can$stress
```

```
[1] 21.58
```

```
n <- rw.25.nmDS.can
```

```
par(mfrow=c(1,1))  
par(mar=c(2,2,1,1), las=0)  
plot(n$points, pch=16, cex=3, ann=FALSE, xaxt='n', yaxt='n',  
      col=ifelse(map$location == 'phone', rgb(0,0,0, .5),  
                 ifelse(map$location == 'index',  
                        rgb(100/255, 149/255, 237/255, .8), # cornflowerblue  
                        rgb(255/255, 140/255, 0/255, .8)))) # darkorange  
mtext('NMDS 1', side=1, line=.0, col='gray40', adj=1)  
mtext('NMDS 2', side=2, line=.0, col='gray40', adj=1)  
legend('topright', legend=c('index', 'thumb', 'phone'),  
       pch=16, col=c('cornflowerblue', 'darkorange', 'gray40'),  
       bty='n', pt.cex=1.5)
```

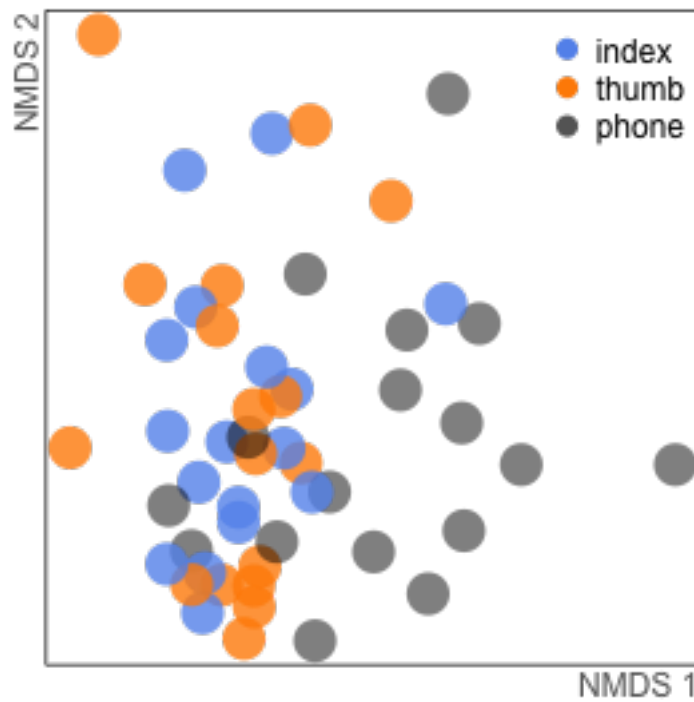


Figure 7: plot of chunk plotNMDS

Earlier, we created sets of variables that combine sample location with gender - this now allows the use of ellipses to visualize confidence intervals in the ordination.

It would be ideal if we could get most information only from sampling a single finger in future cell phone monitoring studies, instead of looking at both index fingers and thumbs. This ordination displays both fingers and phones.

```
# pdf('ordinationGenderBothFingers.pdf', height=4, width=8)
par(mfrow=c(1,2))
par(mar=c(2,2,2,1), las=0)
plot(n$points, type='n', ann=FALSE, xaxt='n', yaxt='n')
mtext('NMDS 1', side=1, line=.3, col='gray40', adj=1)
mtext('NMDS 2', side=2, line=.0, col='gray40', adj=1)
points(n$points[intersect(m,finger), ], pch=21, cex=2, col='gray', bg='white')
points(n$points[intersect(m,p), ], pch=21, cex=2, col='gray', bg='white')
ordiellipse(n, groups=map$loc.gen, show.groups='phone.f',
            draw='polygon', col='gray80')
ordiellipse(n, groups=map$loc.gen, show.groups='phone.f',
            draw='lines', lwd=1.2, col='gray80')
ordiellipse(n, groups=map$loc.gen2, show.groups='finger.f',
            draw='polygon', col='cornflowerblue')
ordiellipse(n, groups=map$loc.gen2, show.groups='finger.f',
            draw='lines', lwd=1.2, col='gray80')
points(n$points[intersect(f,finger), ], pch=21, cex=2.5, col='gray30',
       bg=rgb(100/255, 149/255, 237/255, .8)) # cornflowerblue
points(n$points[intersect(f,p), ], pch=21, cex=2.5, col='gray30',
       bg=rgb(0,0,0, .5))
```

```

mtext('(a) Females', line=.2, font=2, cex=1.5, adj=0)

par(mar=c(2,1,2,2), las=0)
plot(n$points, type='n', ann=FALSE, xaxt='n', yaxt='n')
mtext('NMDS 1', side=1, line=.3, col='gray40', adj=1)
#mtext('NMDS 2', side=2, line=.0, col='gray40', adj=1)
points(n$points[intersect(f,finger), ], pch=21, cex=2, col='gray', bg='white')
points(n$points[intersect(f,p), ], pch=21, cex=2, col='gray', bg='white')
ordiellipse(n, groups=map$loc.gen, show.groups='phone.m',
            draw='polygon', col='gray80')
ordiellipse(n, groups=map$loc.gen, show.groups='phone.m',
            draw='lines', lwd=1.2, col='gray80')
ordiellipse(n, groups=map$loc.gen2, show.groups='finger.m',
            draw='polygon', col='cornflowerblue')
ordiellipse(n, groups=map$loc.gen2, show.groups='finger.m',
            draw='lines', lwd=1.2, col='gray80')
points(n$points[intersect(m,finger), ], pch=21, cex=2.5, col='gray30',
       bg=rgb(100/255, 149/255, 237/255, .8)) # cornflowerblue
points(n$points[intersect(m,p), ], pch=21, cex=2.5, col='gray30',
       bg=rgb(0,0,0, .5))
# text(5,5,'Males', font=2, cex=2, col='gray40')
legend('bottomright', legend=c('phone', 'finger'), pch=21, bty='n', y.intersp=.9,
      pt.bg=c('gray40', 'cornflowerblue'), col='gray30', cex=1, pt.cex=1.5)
mtext('(b) Males', line=.2, font=2, cex=1.5, adj=0)

```

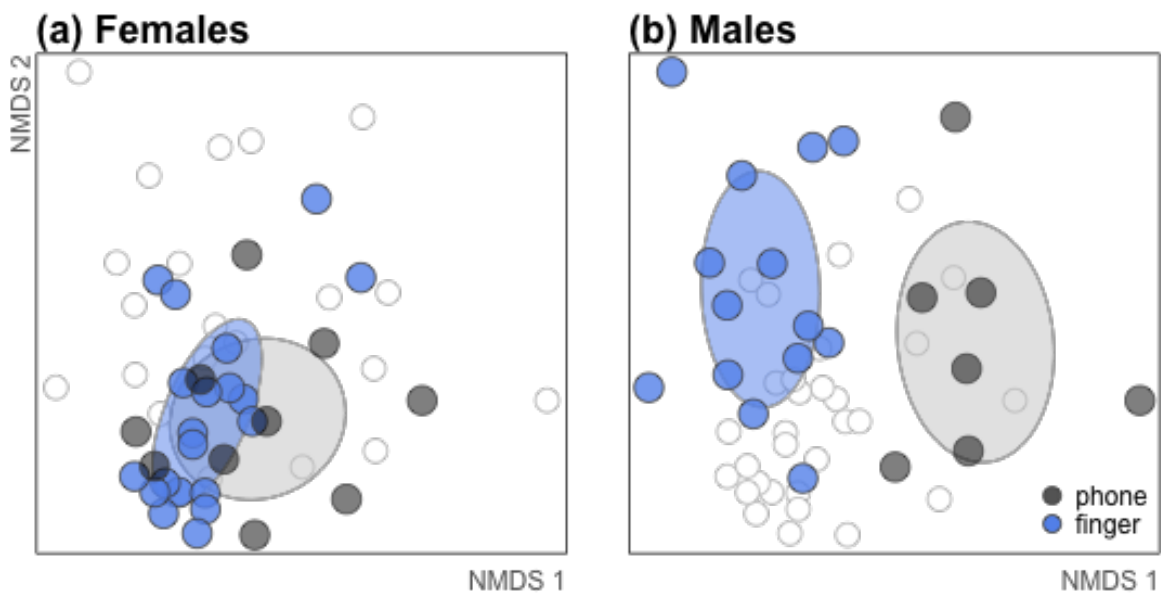


Figure 8: plot of chunk genderFingerNMDS

```
# dev.off()
```

This version uses only index fingers compared to phones.

```

pdf('ordinationGenderIndex.pdf', height=4, width=8)

par(mfrow=c(1,2))
par(mar=c(2,2,2,1), las=0)
plot(n$points, type='n', ann=FALSE, xaxt='n', yaxt='n')
mtext('NMDS 1', side=1, line=.3, col='gray40', adj=1)
mtext('NMDS 2', side=2, line=.0, col='gray40', adj=1)
points(n$points[intersect(m,index), ], pch=21, cex=2, col='gray', bg='white')
points(n$points[intersect(m,p), ], pch=21, cex=2, col='gray', bg='white')
ordiellipse(n, groups=map$loc.gen, show.groups='phone.f',
  draw='polygon', col='gray80')
ordiellipse(n, groups=map$loc.gen, show.groups='phone.f',
  draw='lines', lwd=1.2, col='gray80')
ordiellipse(n, groups=map$loc.gen, show.groups='index.f',
  draw='polygon', col='cornflowerblue')
ordiellipse(n, groups=map$loc.gen, show.groups='index.f',
  draw='lines', lwd=1.2, col='gray80')
points(n$points[intersect(f,index), ], pch=21, cex=2.5, col='gray30',
  bg=rgb(100/255, 149/255, 237/255, .8)) # cornflowerblue
points(n$points[intersect(f,p), ], pch=21, cex=2.5, col='gray30',
  bg=rgb(0,0,0, .5))
mtext('(a) Females', line=.2, font=2, cex=1.5, adj=0)

par(mar=c(2,1,2,2), las=0)
plot(n$points, type='n', ann=FALSE, xaxt='n', yaxt='n')
mtext('NMDS 1', side=1, line=.3, col='gray40', adj=1)
points(n$points[intersect(f,index), ], pch=21, cex=2, col='gray', bg='white')
points(n$points[intersect(f,p), ], pch=21, cex=2, col='gray', bg='white')
ordiellipse(n, groups=map$loc.gen, show.groups='phone.m',
  draw='polygon', col='gray80')
ordiellipse(n, groups=map$loc.gen, show.groups='phone.m',
  draw='lines', lwd=1.2, col='gray80')
ordiellipse(n, groups=map$loc.gen, show.groups='index.m',
  draw='polygon', col='cornflowerblue')
ordiellipse(n, groups=map$loc.gen, show.groups='index.m',
  draw='lines', lwd=1.2, col='gray80')
points(n$points[intersect(m,index), ], pch=21, cex=2.5, col='gray30',
  bg=rgb(100/255, 149/255, 237/255, .8)) # cornflowerblue
points(n$points[intersect(m,p), ], pch=21, cex=2.5, col='gray30',
  bg=rgb(0,0,0, .5))
# text(5,5,'Males', font=2, cex=2, col='gray40')
legend('bottomright', legend=c('phone', 'index'), pch=21, bty='n', y.intersp=.9,
  pt.bg=c('gray40', 'cornflowerblue'), col='gray30', cex=1, pt.cex=1.5)
mtext('(b) Males', line=.2, font=2, cex=1.5, adj=0)

```

```
#dev.off()
```

It seems clear that women and men are falling out in different parts of the ordination. And in fact the difference is highly significant for both phones and fingers - gender makes a difference.

```

can.phones <- as.dist(as.matrix(rw.25.can)[p, p])
can.fingers <- as.dist(as.matrix(rw.25.can)[finger, finger])

```

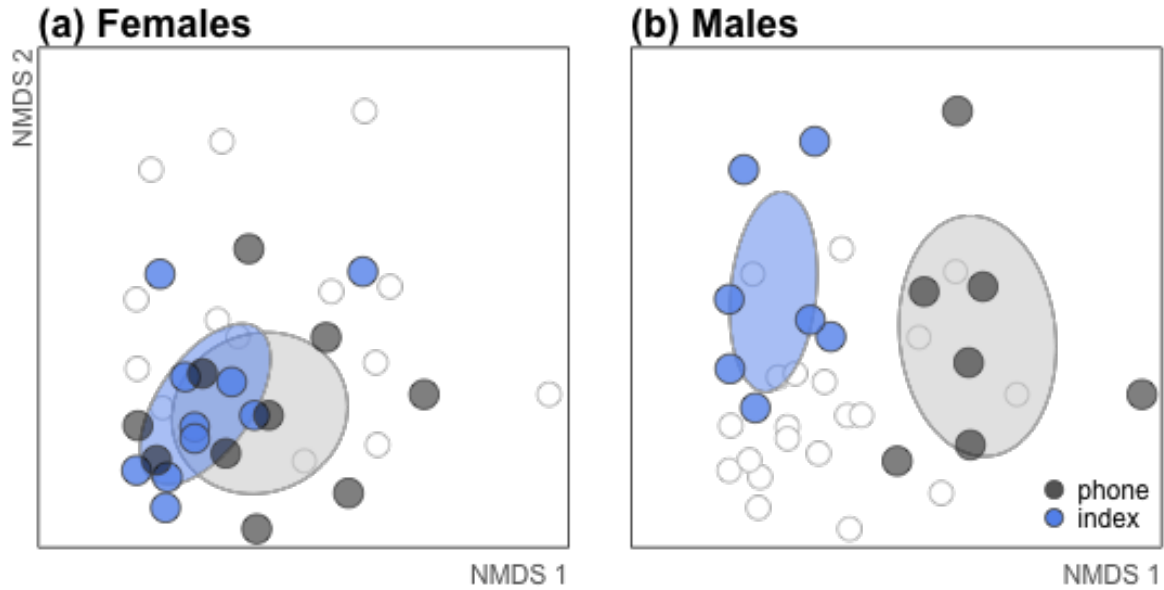


Figure 9: plot of chunk genderIndexNMDS

```
can.index <- as.dist(as.matrix(rw.25.can)[index, thumb])
can.thumb <- as.dist(as.matrix(rw.25.can)[thumb, thumb])

map.phones <- map[p, ]
map.fingers <- map[finger, ]
map.index <- map[index, ]
map.thumb <- map[thumb, ]

adonisAllGender <- adonis(rw.25.can ~ map$gender)$aov.tab
adonisPhoneGender <- adonis(can.phones ~ map.phones$gender)$aov.tab
adonisFingerGender <- adonis(can.fingers ~ map.fingers$gender)$aov.tab
adonisIndexGender <- adonis(can.index ~ map.index$gender)$aov.tab
adonisThumbGender <- adonis(can.thumb ~ map.thumb$gender)$aov.tab

#print(xtable(adonisAllGender))#, type='html')
phtable(adonisAllGender, capt='Gender difference for all samples together?')
```

	Df	SumsOfSqs	MeanSqs	F.Model	R2	Pr(>F)
map\$gender	1	0.65	0.65	1.52	0.03	0.0010
Residuals	49	20.79	0.42		0.97	
Total	50	21.44			1.00	

Table 6: Gender difference for all samples together?

```
phtable(adonisFingerGender, capt='Gender difference for both fingers together?')
```

```
phtable(adonisPhoneGender, capt='Gender difference for just phones?')
```

	Df	SumsOfSqs	MeanSqs	F.Model	R2	Pr(>F)
map.fingers\$gender	1	0.63	0.63	1.51	0.04	0.0010
Residuals	32	13.28	0.42		0.96	
Total	33	13.91			1.00	

Table 7: Gender difference for both fingers together?

	Df	SumsOfSqs	MeanSqs	F.Model	R2	Pr(>F)
map.phones\$gender	1	0.48	0.48	1.11	0.07	0.0110
Residuals	15	6.49	0.43		0.93	
Total	16	6.97			1.00	

Table 8: Gender difference for just phones?

```
phtable(adonisIndexGender, capt='Gender difference for just index fingers?')
```

	Df	SumsOfSqs	MeanSqs	F.Model	R2	Pr(>F)
map.index\$gender	1	0.50	0.50	1.18	0.07	0.0160
Residuals	15	6.31	0.42		0.93	
Total	16	6.81			1.00	

Table 9: Gender difference for just index fingers?

```
phtable(adonisThumbGender, capt='Gender difference for just thumbs?')
```


	Df	SumsOfSqs	MeanSqs	F.Model	R2	Pr(>F)
map.thumb\$gender	1	0.48	0.48	1.14	0.07	0.0280
Residuals	15	6.35	0.42		0.93	
Total	16	6.83			1.00	

Table 10: Gender difference for just thumbs?

Phones, maybe - but also underpowered. Fingers, definitely, but this includes all fingers together which is pseudoreplication. Best to rely on just one finger.

Do index fingers tell the whole story? Or are both fingers together more powerful? It is a tough question because it is an unbalanced comparison, but the nmDS above makes it clear that the significant separation from phones is evident even when only using an index and not grouping both fingers.

```
can.index.phone.f <- as.dist(as.matrix(rw.25.can)[intersect(c(index,p),f),
                                                         intersect(c(index,p),f)])
map.index.phone.f <- map[intersect(c(index,p),f), ]
can.index.phone.m <- as.dist(as.matrix(rw.25.can)[intersect(c(index,p),m),
                                                         intersect(c(index,p),m)])
map.index.phone.m <- map[intersect(c(index,p),m), ]

can.thumb.phone.f <- as.dist(as.matrix(rw.25.can)[intersect(c(thumb,p),f),
                                                         intersect(c(thumb,p),f)])
map.thumb.phone.f <- map[intersect(c(index,p),f), ]
can.thumb.phone.m <- as.dist(as.matrix(rw.25.can)[intersect(c(thumb,p),m),
                                                         intersect(c(thumb,p),m)])
map.thumb.phone.m <- map[intersect(c(thumb,p),m), ]

can.finger.phone.f <- as.dist(as.matrix(rw.25.can)[f, f])
map.finger.phone.f <- map[f, ]
can.finger.phone.m <- as.dist(as.matrix(rw.25.can)[m, m])
map.finger.phone.m <- map[m, ]

adonisIndexPhoneF <- adonis(can.index.phone.f ~ map.index.phone.f$location)$aov.tab
adonisIndexPhoneM <- adonis(can.index.phone.m ~ map.index.phone.m$location)$aov.tab
adonisThumbPhoneF <- adonis(can.thumb.phone.f ~ map.thumb.phone.f$location)$aov.tab
adonisThumbPhoneM <- adonis(can.thumb.phone.m ~ map.thumb.phone.m$location)$aov.tab
adonisFingerPhoneF <- adonis(can.finger.phone.f ~ map.finger.phone.f$location2)$aov.tab
adonisFingerPhoneM <- adonis(can.finger.phone.m ~ map.finger.phone.m$location2)$aov.tab

phtable(adonisIndexPhoneF, capt="Women's phones different from index fingers?")
```

	Df	SumsOfSqs	MeanSqs	F.Model	R2	Pr(>F)
map.index.phone.f\$location	1	0.42	0.42	1.01	0.05	0.3900
Residuals	18	7.57	0.42		0.95	
Total	19	8.00			1.00	

Table 11: Women's phones different from index fingers?

```
phtable(adonisIndexPhoneM, capt="Men's phones different from index fingers?")
```

	Df	SumsOfSqs	MeanSqs	F.Model	R2	Pr(>F)
map.index.phone.m\$location	1	0.51	0.51	1.18	0.09	0.0030
Residuals	12	5.16	0.43		0.91	
Total	13	5.67			1.00	

Table 12: Men's phones different from index fingers?

```
phtable(adonisThumbPhoneF, capt="Women's phones different from thumbs?")
```

	Df	SumsOfSqs	MeanSqs	F.Model	R2	Pr(>F)
map.thumb.phone.f\$location	1	0.42	0.42	1.01	0.05	0.4140
Residuals	18	7.58	0.42		0.95	
Total	19	8.00			1.00	

Table 13: Women's phones different from thumbs?

```
phtable(adonisThumbPhoneM, capt="Men's phones different from thumbs?")
```

	Df	SumsOfSqs	MeanSqs	F.Model	R2	Pr(>F)
map.thumb.phone.m\$location	1	0.52	0.52	1.20	0.09	0.0010
Residuals	12	5.26	0.44		0.91	
Total	13	5.78			1.00	

Table 14: Men's phones different from thumbs?

```
phtable(adonisFingerPhoneF, capt="Women's phones different from both fingers together?")
```

	Df	SumsOfSqs	MeanSqs	F.Model	R2	Pr(>F)
map.finger.phone.f\$location2	1	0.45	0.45	1.08	0.04	0.0970
Residuals	28	11.65	0.42		0.96	
Total	29	12.10			1.00	

Table 15: Women's phones different from both fingers together?

```
phtable(adonisFingerPhoneM, capt="Men's phones different from both fingers together?")
```

	Df	SumsOfSqs	MeanSqs	F.Model	R2	Pr(>F)
map.finger.phone.m\$location2	1	0.57	0.57	1.34	0.07	0.0010
Residuals	19	8.12	0.43		0.93	
Total	20	8.70			1.00	

Table 16: Men's phones different from both fingers together?

Index fingers alone seem to explain the most variation. R^2 is higher, though the F values are slightly lower due to less statistical power. This indicates that index alone is a good finger to use in further studies. Additionally, women don't show a significant difference from the communities on their phones, while men do! Perhaps women will be easier to track by their phones? This contrasts slightly from the results above that say that both men and women share more OTUs with their phones than other people's phones. So this indicates that whether the phone can be used to track person microbiomes will depend on how we look at the data.

```
ls()
```

```
[1] "adonisAllGender"      "adonisFingerGender" "adonisFingerPhoneF"
[4] "adonisFingerPhoneM"  "adonisIndexGender"  "adonisIndexPhoneF"
[7] "adonisIndexPhoneM"   "adonisPhoneGender"  "adonisThumbGender"
[10] "adonisThumbPhoneF"   "adonisThumbPhoneM"  "bar.df"
[13] "bar.df.female.j"     "bar.df.male.j"       "bar.df.nowash.j"
[16] "bar.df.wash.j"       "bar.df2"             "bar.female.j"
[19] "bar.jac"             "bar.jac.df"          "bar.male.j"
[22] "bar.mf"              "bar.nowash.j"        "bar.others"
[25] "bar.others.df"       "bar.others.f.df"     "bar.others.m.df"
[28] "bar.summary.f"       "bar.summary.m"       "bar.summary2"
[31] "bar.wash.j"          "bar.wnw"             "barname"
[34] "can.finger.phone.f"  "can.finger.phone.m"  "can.fingers"
[37] "can.index"           "can.index.phone.f"   "can.index.phone.m"
[40] "can.phones"          "can.thumb"           "can.thumb.phone.f"
[43] "can.thumb.phone.m"   "diff.f"              "diff.m"
[46] "dis"                 "dismat"              "Evenness"
[49] "f"                   "finger"              "fir.mean"
[52] "fir.se"              "fir.table"           "fir.table.10"
[55] "fir.taxo"            "fir.taxo.10"         "fixit"
[58] "fixx"                "fixy"                "font"
[61] "i"                   "index"               "JacJFM"
[64] "m"                   "makeBarDF"           "makeBarSummary"
[67] "makeTaxo"            "map"                 "map.finger.phone.f"
[70] "map.finger.phone.m"  "map.fingers"         "map.index"
[73] "map.index.phone.f"   "map.index.phone.m"   "map.phones"
[76] "map.thumb"           "map.thumb.phone.f"   "map.thumb.phone.m"
[79] "medAbu5"             "medAbu5.f"           "medAbu5.fing"
[82] "medAbu5.m"           "medAbu5.p"           "mids"
[85] "n"                   "p"                   "ph"
[88] "ph.mean"             "ph.se"               "pro.mean"
[91] "pro.se"              "pro.table"           "pro.table.10"
[94] "pro.taxo"            "pro.taxo.10"         "phtable"
[97] "QiimeIn"             "realSE"              "rw.25"
[100] "rw.25.can"           "rw.25.nmnds.can"     "rw.25.rel"
[103] "rw.taxo.25"          "se"                  "taxo"
[106] "test"                "thumb"               "ylim"
```

```
save.image('phones_knitr.RData')
```

This cluttered workspace *should* have everything necessary to reproduce any analysis or figure separately without running the entire script.